

ProjectPhase3

1. Model: Logistic Regression

```
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr     1.1.4     v readr     2.1.5
v forcats   1.0.0     v stringr   1.5.2
v ggplot2   4.0.0     v tibble    3.3.0
v lubridate  1.9.4     v tidyr    1.3.1
v purrr     1.1.0
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()    masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to becom
```

```
#Import Data
bank_data<- read.csv2("bank-full.csv")
```

```
str(bank_data)
```

```
'data.frame': 45211 obs. of 17 variables:
 $ age      : int 58 44 33 47 33 35 28 42 58 43 ...
 $ job      : chr "management" "technician" "entrepreneur" "blue-collar" ...
 $ marital   : chr "married" "single" "married" "married" ...
 $ education: chr "tertiary" "secondary" "secondary" "unknown" ...
 $ default   : chr "no" "no" "no" "no" ...
 $ balance   : int 2143 29 2 1506 1 231 447 2 121 593 ...
 $ housing   : chr "yes" "yes" "yes" "yes" ...
 $ loan      : chr "no" "no" "yes" "no" ...
```

```

$ contact  : chr  "unknown" "unknown" "unknown" "unknown" ...
$ day      : int  5 5 5 5 5 5 5 5 5 ...
$ month    : chr  "may" "may" "may" "may" ...
$ duration : int  261 151 76 92 198 139 217 380 50 55 ...
$ campaign : int  1 1 1 1 1 1 1 1 1 ...
$ pdays    : int  -1 -1 -1 -1 -1 -1 -1 -1 -1 ...
$ previous : int  0 0 0 0 0 0 0 0 0 ...
$ poutcome : chr  "unknown" "unknown" "unknown" "unknown" ...
$ y        : chr  "no" "no" "no" "no" ...

```

```
summary(bank_data)
```

age	job	marital	education
Min. :18.00	Length:45211	Length:45211	Length:45211
1st Qu.:33.00	Class :character	Class :character	Class :character
Median :39.00	Mode :character	Mode :character	Mode :character
Mean :40.94			
3rd Qu.:48.00			
Max. :95.00			
default	balance	housing	loan
Length:45211	Min. : -8019	Length:45211	Length:45211
Class :character	1st Qu.: 72	Class :character	Class :character
Mode :character	Median : 448	Mode :character	Mode :character
	Mean : 1362		
	3rd Qu.: 1428		
	Max. :102127		
contact	day	month	duration
Length:45211	Min. : 1.00	Length:45211	Min. : 0.0
Class :character	1st Qu.: 8.00	Class :character	1st Qu.: 103.0
Mode :character	Median :16.00	Mode :character	Median : 180.0
	Mean :15.81		Mean : 258.2
	3rd Qu.:21.00		3rd Qu.: 319.0
	Max. :31.00		Max. :4918.0
campaign	pdays	previous	poutcome
Min. : 1.000	Min. : -1.0	Min. : 0.0000	Length:45211
1st Qu.: 1.000	1st Qu.: -1.0	1st Qu.: 0.0000	Class :character
Median : 2.000	Median : -1.0	Median : 0.0000	Mode :character
Mean : 2.764	Mean : 40.2	Mean : 0.5803	
3rd Qu.: 3.000	3rd Qu.: -1.0	3rd Qu.: 0.0000	
Max. :63.000	Max. :871.0	Max. :275.0000	
y			
Length:45211			

```

Class :character
Mode  :character

# set the target as factor "y"
bank_data$y <- factor(bank_data$y, levels = c("no", "yes"))

# Remove leakage variable
bank_data$duration <- NULL

# Fix pdays based on your dataset ( -1 = not contacted )
bank_data$previous_contact <- ifelse(bank_data$pdays == -1, "no", "yes")
bank_data$previous_contact <- factor(bank_data$previous_contact)

#Convert categorical variables to factors
factor_vars <- c("job","marital","education","default","housing",
                 "loan","contact","month","poutcome","previous_contact")

bank_data[factor_vars] <- lapply(bank_data[factor_vars], factor)

# Train-test split
set.seed(123) # ensures reproducibility

n <- nrow(bank_data)
train_index <- sample(1:n, size = 0.7 * n)

train <- bank_data[train_index, ]
test  <- bank_data[-train_index, ]

# Check dimensions
dim(train)

```

[1] 31647 17

```
dim(test)
```

[1] 13564 17

```
prop.table(table(train$y))
```

```
no      yes  
0.8823585 0.1176415
```

```
prop.table(table(test$y))
```

```
no      yes  
0.8845473 0.1154527
```

```
#Logistic Regression
```

```
full_model <- glm(  
  y ~ age + job + marital + education + default + housing +  
    loan + contact + month + campaign + previous + previous_contact + poutcome + balance,  
  data = train,  
  family = binomial  
)  
  
summary(full_model)
```

Call:

```
glm(formula = y ~ age + job + marital + education + default +  
  housing + loan + contact + month + campaign + previous +  
  previous_contact + poutcome + balance, family = binomial,  
  data = train)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-3.298e+00	1.449e+00	-2.276	0.022854 *
age	-1.185e-03	2.343e-03	-0.506	0.613074
jobblue-collar	-1.128e-01	7.708e-02	-1.464	0.143240
jobentrepreneur	-1.057e-01	1.304e-01	-0.810	0.417764
jobhousemaid	-2.950e-01	1.418e-01	-2.080	0.037567 *
jobmanagement	-4.392e-02	7.842e-02	-0.560	0.575409
jobretired	4.246e-01	1.044e-01	4.068	4.75e-05 ***
jobsself-employed	-1.038e-01	1.176e-01	-0.882	0.377550

jobservices	-1.210e-01	8.922e-02	-1.356	0.175031
jobstudent	2.646e-01	1.201e-01	2.203	0.027625 *
jobtechnician	-8.526e-02	7.367e-02	-1.157	0.247140
jobunemployed	1.559e-01	1.152e-01	1.353	0.176137
jobunknown	-3.275e-01	2.656e-01	-1.233	0.217525
maritalmarried	-2.257e-01	6.154e-02	-3.667	0.000245 ***
maritalsingle	5.682e-02	7.045e-02	0.807	0.419899
educationsecondary	1.376e-01	6.764e-02	2.034	0.041986 *
educationtertiary	2.821e-01	7.888e-02	3.577	0.000348 ***
educationunknown	1.459e-01	1.116e-01	1.307	0.191130
defaultyes	-1.520e-01	1.744e-01	-0.872	0.383229
housingyes	-5.302e-01	4.559e-02	-11.629	< 2e-16 ***
loanyes	-3.765e-01	6.293e-02	-5.983	2.19e-09 ***
contacttelephone	-2.558e-01	7.879e-02	-3.247	0.001166 **
contactunknown	-1.398e+00	7.579e-02	-18.442	< 2e-16 ***
monthaug	-9.783e-01	8.237e-02	-11.877	< 2e-16 ***
monthdec	5.188e-01	1.930e-01	2.688	0.007185 **
monthfeb	-4.799e-01	8.875e-02	-5.407	6.41e-08 ***
monthjan	-1.147e+00	1.253e-01	-9.150	< 2e-16 ***
monthjul	-7.504e-01	7.960e-02	-9.427	< 2e-16 ***
monthjun	9.130e-02	9.565e-02	0.955	0.339824
monthmar	9.311e-01	1.343e-01	6.935	4.06e-12 ***
monthmay	-5.415e-01	7.461e-02	-7.257	3.94e-13 ***
monthnov	-9.921e-01	8.887e-02	-11.163	< 2e-16 ***
monthoct	5.305e-01	1.168e-01	4.542	5.56e-06 ***
monthsep	6.696e-01	1.269e-01	5.278	1.30e-07 ***
campaign	-7.487e-02	9.709e-03	-7.711	1.25e-14 ***
previous	8.086e-03	6.215e-03	1.301	0.193279
previous_contactyes	2.239e+00	1.442e+00	1.553	0.120402
poutcomeother	3.049e-01	9.492e-02	3.212	0.001318 **
poutcomesuccess	2.276e+00	8.787e-02	25.900	< 2e-16 ***
poutcomeunknown	2.375e+00	1.443e+00	1.646	0.099776 .
balance	2.558e-05	5.476e-06	4.670	3.01e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 22925 on 31646 degrees of freedom
 Residual deviance: 19071 on 31606 degrees of freedom
 AIC: 19153

Number of Fisher Scoring iterations: 6

```
step_model <- step(full_model, direction = "both", trace = FALSE)
summary(step_model)
```

Call:

```
glm(formula = y ~ job + marital + education + housing + loan +
    contact + month + campaign + previous_contact + poutcome +
    balance, family = binomial, data = train)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-3.350e+00	1.444e+00	-2.320	0.020357 *
jobblue-collar	-1.134e-01	7.702e-02	-1.473	0.140808
jobentrepreneur	-1.108e-01	1.303e-01	-0.850	0.395382
jobhousemaid	-3.014e-01	1.415e-01	-2.130	0.033177 *
jobmanagement	-4.628e-02	7.832e-02	-0.591	0.554561
jobretired	4.004e-01	9.424e-02	4.249	2.14e-05 ***
jobsself-employed	-1.073e-01	1.175e-01	-0.913	0.361377
jobservices	-1.209e-01	8.918e-02	-1.356	0.175202
jobstudent	2.762e-01	1.181e-01	2.339	0.019322 *
jobtechnician	-8.580e-02	7.366e-02	-1.165	0.244053
jobunemployed	1.538e-01	1.152e-01	1.336	0.181685
jobunknown	-3.333e-01	2.655e-01	-1.255	0.209344
maritalmarried	-2.217e-01	6.131e-02	-3.616	0.000299 ***
maritalsingle	6.935e-02	6.607e-02	1.050	0.293862
educationsecondary	1.404e-01	6.724e-02	2.089	0.036747 *
educationtertiary	2.878e-01	7.817e-02	3.681	0.000232 ***
educationunknown	1.444e-01	1.116e-01	1.294	0.195619
housingyes	-5.266e-01	4.538e-02	-11.605	< 2e-16 ***
loanyes	-3.795e-01	6.279e-02	-6.045	1.50e-09 ***
contacttelephone	-2.595e-01	7.805e-02	-3.325	0.000883 ***
contactunknown	-1.400e+00	7.576e-02	-18.481	< 2e-16 ***
monthaug	-9.796e-01	8.231e-02	-11.902	< 2e-16 ***
monthdec	5.206e-01	1.930e-01	2.698	0.006984 **
monthfeb	-4.782e-01	8.868e-02	-5.393	6.93e-08 ***
monthjan	-1.146e+00	1.254e-01	-9.141	< 2e-16 ***
monthjul	-7.516e-01	7.954e-02	-9.450	< 2e-16 ***
monthjun	9.272e-02	9.565e-02	0.969	0.332360
monthmar	9.316e-01	1.342e-01	6.941	3.90e-12 ***
monthmay	-5.405e-01	7.457e-02	-7.248	4.24e-13 ***
monthnov	-9.944e-01	8.880e-02	-11.198	< 2e-16 ***
monthoct	5.306e-01	1.167e-01	4.546	5.46e-06 ***

```

monthsep          6.710e-01  1.268e-01  5.291  1.22e-07 ***
campaign         -7.458e-02  9.696e-03  -7.692  1.45e-14 ***
previous_contactyes 2.255e+00  1.441e+00  1.565  0.117619
poutcomeother    3.143e-01  9.459e-02  3.323  0.000890 ***
poutcomesuccess  2.278e+00  8.785e-02  25.930 < 2e-16 ***
poutcomeunknown  2.367e+00  1.442e+00  1.642  0.100657
balance          2.558e-05  5.452e-06  4.692  2.70e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

(Dispersion parameter for binomial family taken to be 1)

```

Null deviance: 22925  on 31646  degrees of freedom
Residual deviance: 19074  on 31609  degrees of freedom
AIC: 19150

```

Number of Fisher Scoring iterations: 6

```
library(car)
```

Loading required package: carData

Attaching package: 'car'

The following object is masked from 'package:dplyr':

```
recode
```

The following object is masked from 'package:purrr':

```
some
```

```
vif(step_model)
```

	GVIF	Df	GVIF^(1/(2*Df))
job	2.887692	11	1.049383
marital	1.168093	2	1.039607
education	2.219748	3	1.142135
housing	1.369677	1	1.170332

loan	1.057187	1	1.028196
contact	1.801797	2	1.158581
month	2.703957	11	1.046252
campaign	1.069901	1	1.034360
previous_contact	1111.354867	1	33.336989
poutcome	1149.581011	3	3.236607
balance	1.036594	1	1.018133

```
final_model <- glm(
  y ~ job + marital + education + housing + loan + contact +
  month + campaign + poutcome + balance,
  data = train,
  family = binomial
)

summary(final_model)
```

Call:

```
glm(formula = y ~ job + marital + education + housing + loan +
  contact + month + campaign + poutcome + balance, family = binomial,
  data = train)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.098e+00	1.278e-01	-8.587	< 2e-16 ***
jobblue-collar	-1.133e-01	7.702e-02	-1.471	0.141258
jobentrepreneur	-1.109e-01	1.303e-01	-0.851	0.394807
jobhousemaid	-3.009e-01	1.415e-01	-2.126	0.033482 *
jobmanagement	-4.639e-02	7.832e-02	-0.592	0.553642
jobretired	4.029e-01	9.422e-02	4.276	1.90e-05 ***
jobself-employed	-1.072e-01	1.175e-01	-0.912	0.361683
jobservices	-1.209e-01	8.918e-02	-1.356	0.175082
jobstudent	2.765e-01	1.181e-01	2.342	0.019166 *
jobtechnician	-8.491e-02	7.365e-02	-1.153	0.248938
jobunemployed	1.543e-01	1.152e-01	1.339	0.180475
jobunknown	-3.327e-01	2.655e-01	-1.253	0.210118
maritalmarried	-2.214e-01	6.131e-02	-3.611	0.000305 ***
maritalsingle	7.014e-02	6.607e-02	1.062	0.288412
educationsecondary	1.408e-01	6.724e-02	2.094	0.036259 *
educationtertiary	2.884e-01	7.817e-02	3.689	0.000225 ***
educationunknown	1.448e-01	1.116e-01	1.298	0.194396

```

housingyes      -5.252e-01  4.537e-02 -11.576 < 2e-16 ***
loanyes        -3.766e-01  6.272e-02 -6.006 1.91e-09 ***
contacttelephone -2.599e-01  7.805e-02 -3.329 0.000870 ***
contactunknown   -1.402e+00  7.576e-02 -18.504 < 2e-16 ***
monthhaug       -9.784e-01  8.230e-02 -11.889 < 2e-16 ***
monthdec         5.211e-01  1.930e-01  2.700 0.006924 **
monthfeb        -4.780e-01  8.868e-02 -5.390 7.04e-08 ***
monthjan         -1.146e+00  1.254e-01 -9.138 < 2e-16 ***
monthjul         -7.519e-01  7.953e-02 -9.454 < 2e-16 ***
monthjun         9.579e-02  9.563e-02  1.002 0.316519
monthmar         9.320e-01  1.342e-01  6.944 3.80e-12 ***
monthmay        -5.402e-01  7.457e-02 -7.245 4.33e-13 ***
monthnov         -9.943e-01  8.880e-02 -11.197 < 2e-16 ***
monthoct         5.310e-01  1.167e-01  4.550 5.36e-06 ***
monthsep         6.716e-01  1.268e-01  5.296 1.18e-07 ***
campaign        -7.451e-02  9.692e-03 -7.688 1.49e-14 ***
poutcomeother   3.144e-01  9.459e-02  3.324 0.000888 ***
poutcomesuccess 2.278e+00  8.784e-02  25.932 < 2e-16 ***
poutcomeunknown 1.126e-01  6.161e-02  1.829 0.067470 .
balance          2.560e-05  5.452e-06  4.695 2.67e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

(Dispersion parameter for binomial family taken to be 1)

```

Null deviance: 22925  on 31646  degrees of freedom
Residual deviance: 19076  on 31610  degrees of freedom
AIC: 19150

```

Number of Fisher Scoring iterations: 6

```
vif(final_model)
```

	GVIF	Df	GVIF^(1/(2*Df))
job	2.885854	11	1.049353
marital	1.167934	2	1.039572
education	2.219891	3	1.142147
housing	1.369114	1	1.170091
loan	1.055712	1	1.027479
contact	1.801943	2	1.158605
month	2.702727	11	1.046230
campaign	1.069844	1	1.034332

```
poutcome 1.216169 3          1.033155  
balance   1.036596 1          1.018133
```

```
test$prob <- predict(final_model, newdata = test, type = "response")  
test$pred <- ifelse(test$prob > 0.5, "yes", "no")  
  
#Confusion Metrics  
table(Predicted = test$pred, Actual = test$y)
```

		Actual
Predicted	no	yes
no	11842	1282
yes	156	284

```
#Accuracy  
mean(test$pred == test$y)
```

```
[1] 0.8939841
```

```
# ROC & AUC  
library(pROC)
```

```
Type 'citation("pROC")' for a citation.
```

```
Attaching package: 'pROC'
```

```
The following objects are masked from 'package:stats':
```

```
cov, smooth, var
```

```
roc_obj <- roc(test$y, test$prob)
```

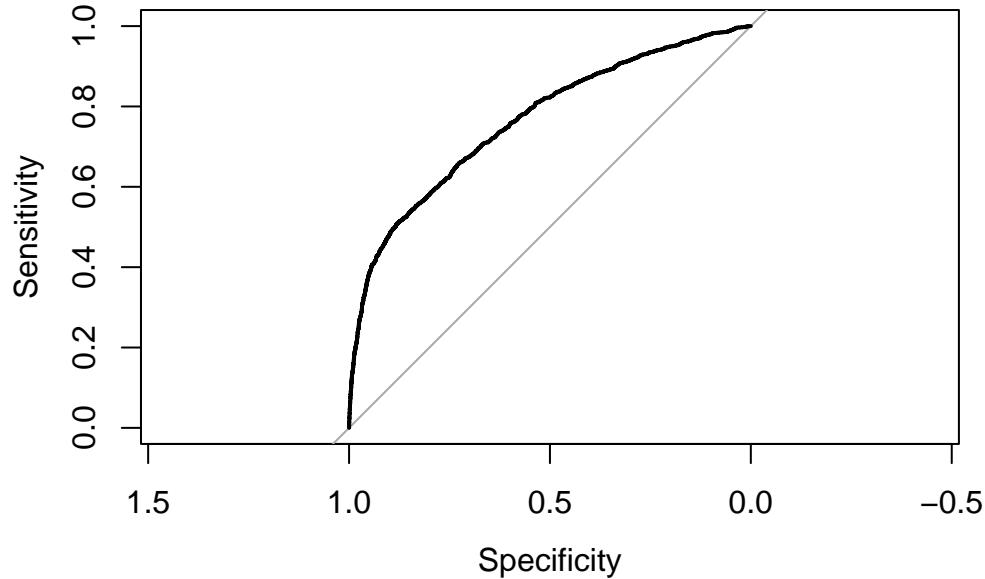
```
Setting levels: control = no, case = yes
```

```
Setting direction: controls < cases
```

```
auc(roc_obj)
```

Area under the curve: 0.7644

```
plot(roc_obj)
```



```
#Try different cutoff

thresholds <- seq(0.1, 0.5, by = 0.05)

for (t in thresholds) {
  pred <- ifelse(test$prob > t, "yes", "no")
  cm <- table(Predicted = pred, Actual = test$y)
  sensitivity <- cm["yes","yes"] / (cm["yes","yes"] + cm["no","yes"])
  specificity <- cm["no","no"] / (cm["no","no"] + cm["yes","no"])
  cat("\nThreshold:", t,
      "Sensitivity:", round(sensitivity,3),
      "Specificity:", round(specificity,3))
}
```

```

Threshold: 0.1 Sensitivity: 0.71 Specificity: 0.659
Threshold: 0.15 Sensitivity: 0.522 Specificity: 0.865
Threshold: 0.2 Sensitivity: 0.434 Specificity: 0.926
Threshold: 0.25 Sensitivity: 0.383 Specificity: 0.95
Threshold: 0.3 Sensitivity: 0.335 Specificity: 0.96
Threshold: 0.35 Sensitivity: 0.29 Specificity: 0.968
Threshold: 0.4 Sensitivity: 0.255 Specificity: 0.976
Threshold: 0.45 Sensitivity: 0.213 Specificity: 0.981
Threshold: 0.5 Sensitivity: 0.181 Specificity: 0.987

```

```
coords(roc_obj, "best", ret="threshold", best.method="youden")
```

```

threshold
1 0.1550712

```

We selected a **0.155** threshold

The default 0.5 threshold produced very low sensitivity because only a small proportion of customers subscribe, causing the model to classify most cases as ‘no.’ To avoid losing high-potential customers, we applied Youden’s Index to identify the optimal decision threshold. The best cutoff was 0.155, which provides the strongest balance between sensitivity and specificity. Lowering the threshold allows the model to capture substantially more potential subscribers while maintaining acceptable precision—making it a more effective strategy for targeted marketing.

```

# Best threshold from Youden Index
best_t <- 0.1550712 # or use: coords(roc_obj, "best", ret="threshold", best.method="youden")

# Predict using the new threshold
test$pred_best <- ifelse(test$prob > best_t, "yes", "no")

# Confusion matrix
cm_best <- table(Predicted = test$pred_best, Actual = test$y)
cm_best

```

		Actual
Predicted	no	yes
no	10509	762
yes	1489	804

```

# Accuracy
accuracy_best <- mean(test$pred_best == test$y)
accuracy_best

[1] 0.834046

# Sensitivity (Recall for "yes")
sensitivity_best <- cm_best["yes", "yes"] /
                      (cm_best["yes", "yes"] + cm_best["no", "yes"])
sensitivity_best

[1] 0.51341

# Specificity (Recall for "no")
specificity_best <- cm_best["no", "no"] /
                      (cm_best["no", "no"] + cm_best["yes", "no"])
specificity_best

[1] 0.875896

# Optional: Print everything clearly
cat("\nThreshold:", best_t,
    "\nAccuracy:", round(accuracy_best,4),
    "\nSensitivity:", round(sensitivity_best,4),
    "\nSpecificity:", round(specificity_best,4), "\n")

```

Threshold: 0.1550712
 Accuracy: 0.834
 Sensitivity: 0.5134
 Specificity: 0.8759

1. Model: Logistic Regression (Summary)

Data Cleaning & Preparation

We converted categorical variables into factors and engineered the `previous_contact` variable to reflect prior outreach more clearly. This ensured that the logistic regression model could interpret the predictors correctly without structural issues.

Logistic Regression Model

The full model allowed us to assess the initial influence of all predictors, but it also revealed several insignificant or overlapping variables, so we used AIC-based stepwise selection to remove redundant or non-informative variables while retaining predictors that meaningfully improved overall model fit. This produced a more stable and interpretable model.

Key Predictors Identified

After stepwise refinement, several predictors remained highly influential. Positive drivers included higher education levels (secondary and tertiary), retired or student job status, successful prior campaign outcomes, and higher account balance. Negative influences included holding housing or personal loans, contacting customers via telephone or unknown channels, and certain campaign months such as July, August, and November. These directional effects provide an initial understanding of the customer segments more likely or less likely to respond.

Model Evaluation (ROC & AUC)

The logistic model achieved an **AUC** of approximately **0.76**, indicating solid baseline discrimination despite the dataset's class imbalance.

Threshold Optimization (Youden Index)

Because the default 0.50 threshold missed many potential subscribers, we applied Youden's Index to find the optimal cutoff. A threshold of **0.155** provided the best balance between sensitivity and specificity, making the model more effective for identifying potential 'yes' cases.

2. Random Forest

```
library(randomForest)
```

```
randomForest 4.7-1.2
```

```
Type rfNews() to see new features/changes/bug fixes.
```

```
Attaching package: 'randomForest'
```

```
The following object is masked from 'package:dplyr':
```

```
combine
```

```
The following object is masked from 'package:ggplot2':
```

```
margin
```

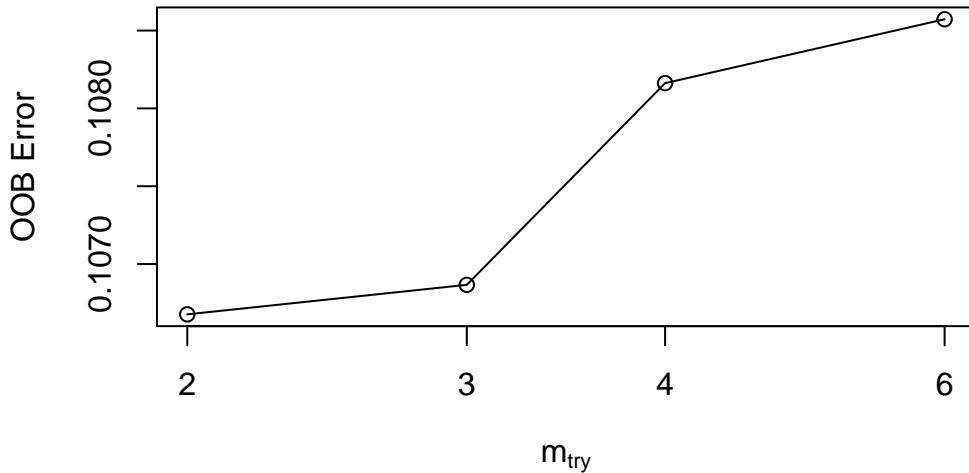
```
set.seed(123)

x_train <- subset(train, select = -y)
y_train <- train$y
```

```
set.seed(123)

tune_res <- tuneRF(
  x = x_train,
  y = y_train,
  stepFactor = 1.5,
  improve = 0.005,
  ntreeTry = 300,
  trace = TRUE,
  plot = TRUE )
```

```
mtry = 4 OOB error = 10.82%
Searching left ...
mtry = 3 OOB error = 10.69%
0.0119778 0.005
mtry = 2 OOB error = 10.67%
0.001774098 0.005
Searching right ...
mtry = 6 OOB error = 10.86%
-0.01596688 0.005
```



```
set.seed(123)

rf_model <- randomForest(
  y ~ .,
  data = train,
  ntree = 500,
  mtry = 2,
  importance = TRUE
)

rf_model
```

Call:

```
randomForest(formula = y ~ ., data = train, ntree = 500, mtry = 2,           importance = TRUE)
  Type of random forest: classification
  Number of trees: 500
  No. of variables tried at each split: 2

  OOB estimate of  error rate: 10.68%
```

Confusion matrix:

	no	yes	class.error
--	----	-----	-------------

```
no  27626 298  0.01067182
yes 3081 642  0.82755842
```

```
rf_prob <- predict(rf_model, newdata=test, type="prob") [,2]
library(pROC)
roc_rf <- roc(test$y, rf_prob)
```

```
Setting levels: control = no, case = yes
```

```
Setting direction: controls < cases
```

```
coords(roc_rf, "best", ret="threshold")
```

```
threshold
1      0.057
```

```
library(pROC)
rf_prob <- predict(rf_model, newdata=test, type="prob") [,2]
roc_rf <- roc(test$y, rf_prob)
```

```
Setting levels: control = no, case = yes
```

```
Setting direction: controls < cases
```

```
best_t <- as.numeric(coords(roc_rf, "best", ret="threshold"))

rf_pred2 <- factor(ifelse(rf_prob > best_t, "yes", "no"),
                     levels = c("no","yes"))

table(Predicted = rf_pred2, Actual = test$y)
```

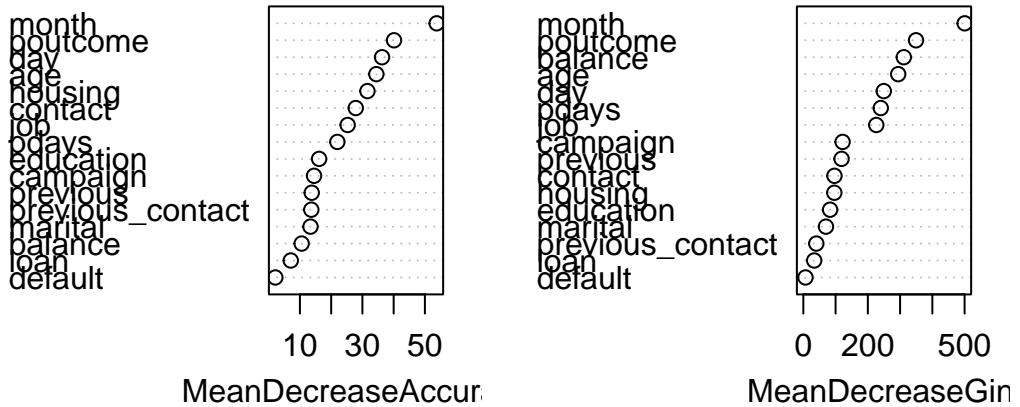
Predicted	Actual	no	yes
no	10099	601	
yes	1899	965	

```
importance(rf_model)
```

	no	yes	MeanDecreaseAccuracy	MeanDecreaseGini
age	31.534066	7.072938	34.464131	294.19204
job	26.859409	-5.871211	25.279882	226.24705
marital	10.757869	5.798110	13.386577	70.18623
education	17.148381	-1.599597	16.146582	83.02434
default	1.006874	3.094812	2.132085	6.64752
balance	7.209322	6.992791	10.557528	311.56701
housing	28.547696	8.445231	31.614042	96.11769
loan	-3.318283	16.033670	7.077996	33.74744
contact	25.621392	12.582831	27.872002	96.54402
day	36.246415	-2.960642	36.284896	249.48974
month	51.445738	12.709346	53.772727	500.32924
campaign	9.622106	11.863759	14.524867	121.75171
pdays	20.227198	14.981491	21.984929	240.12281
previous	12.950615	11.588766	13.780203	118.59273
poutcome	31.941247	2.695573	40.115277	349.41530
previous_contact	13.049454	9.885696	13.646109	40.26791

```
varImpPlot(rf_model,
            sort = TRUE,
            n.var = min(20, ncol(train)-1),
            main = "Random Forest Variable Importance")
```

Random Forest Variable Importance



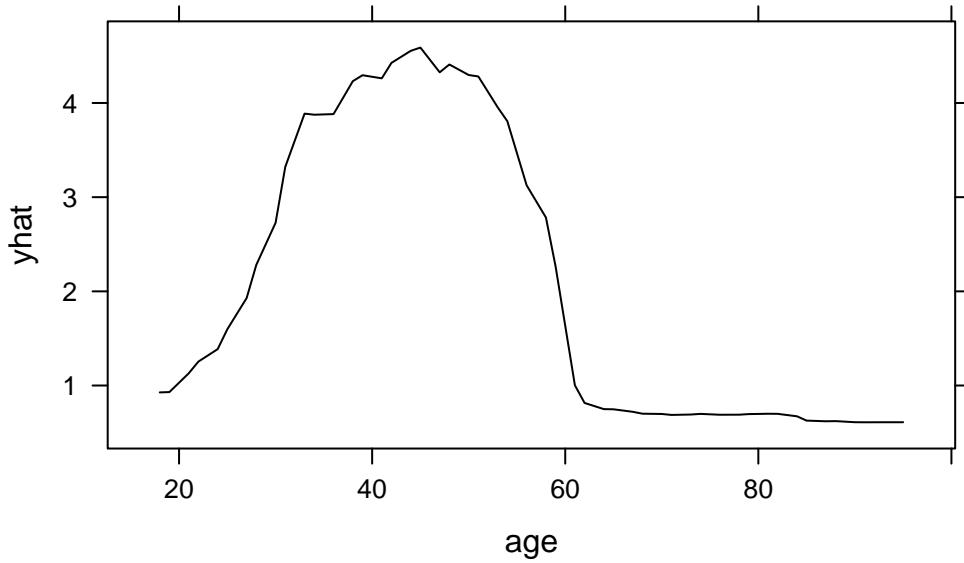
```
library(pdp)
```

```
Attaching package: 'pdp'
```

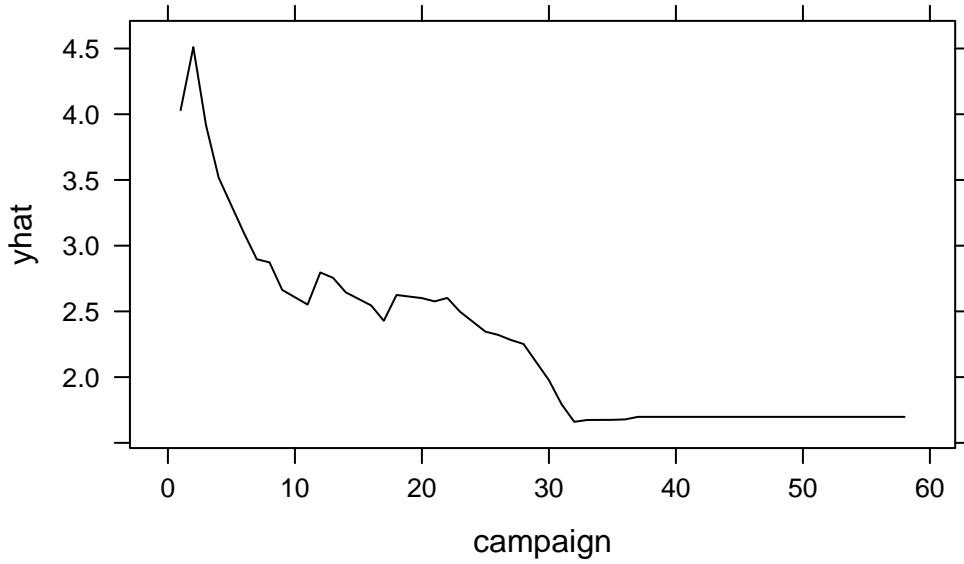
```
The following object is masked from 'package:purrr':
```

```
partial
```

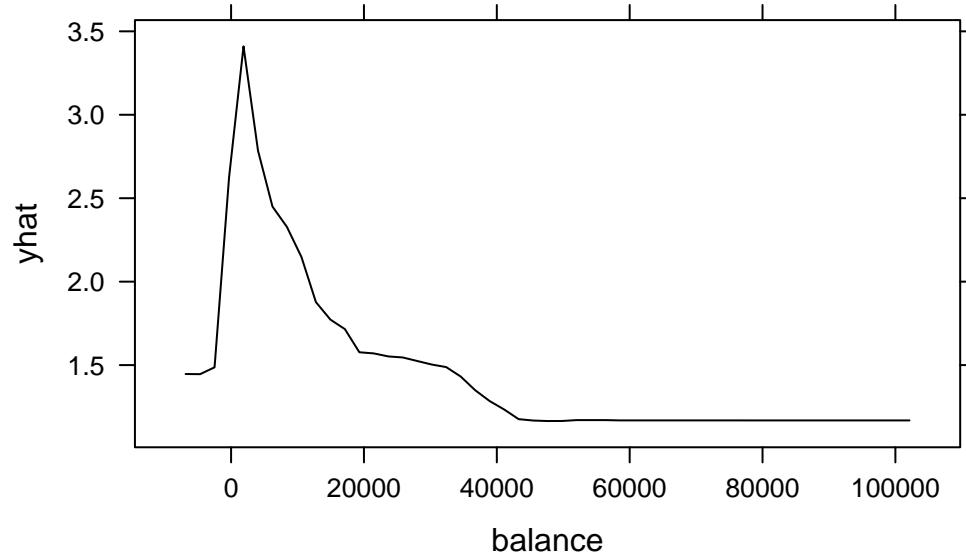
```
partial(rf_model, pred.var = "age", plot = TRUE)
```



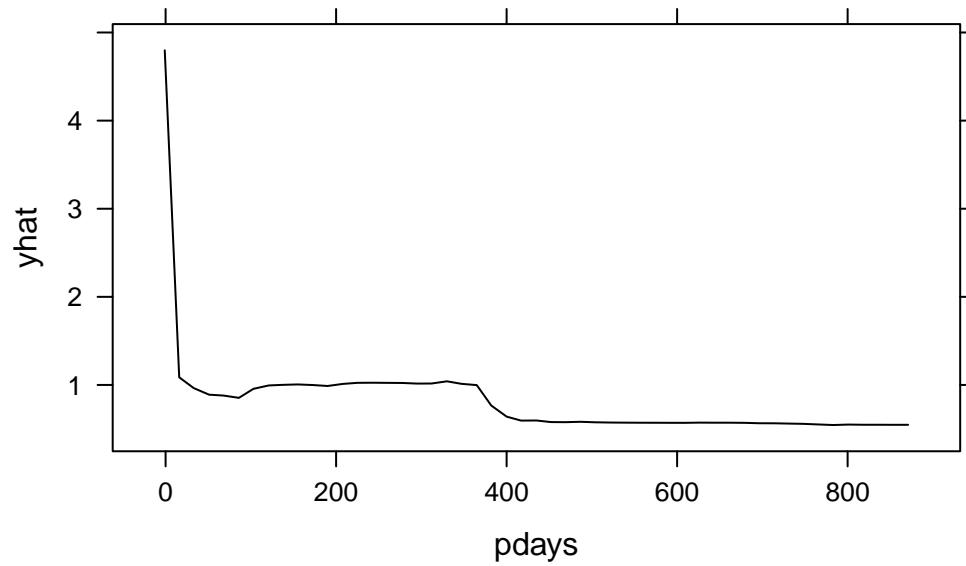
```
partial(rf_model, pred.var = "campaign", plot = TRUE)
```



```
partial(rf_model, pred.var = "balance", plot = TRUE)
```



```
partial(rf_model, pred.var = "pdays", plot = TRUE)
```



2. Random Forest model (Summary)

We trained a Random Forest model (500 trees) to capture nonlinear patterns and identify which variables most strongly influence subscription to the term deposit. We extracted variable importance rankings and generated partial dependence plots for the top predictors.

Key Findings

1. **Month is the strongest predictor:** Subscription probability varies significantly by contact month, indicating clear seasonality effects.
2. **Previous campaign outcome strongly predicts success:** Customers with a prior successful interaction have a much higher likelihood of subscribing.
3. **Age shows a nonlinear pattern:** Middle-aged customers (around 30–55) have the highest predicted probability. Younger and older customers show lower interest.
4. **pdays indicates recency matters:** Customers contacted recently, or first-time contacts, are more likely to subscribe than those contacted a long time ago.
5. **Balance shows diminishing returns:** Moderate balances correspond to higher predicted probability, while very low and very high balances are associated with lower probability.
6. **Campaign frequency has a negative effect:** As the number of contact attempts increases, subscription probability declines, suggesting diminishing effectiveness with repeated calls.

The Random Forest model suggests that campaign timing, prior interactions, age, and contact recency are key drivers of customer response. Targeting should prioritize middle-aged customers during high-performing months, especially those with successful prior outcomes or recent contact

3. Decision Tree

```
library(rpart)
library(rpart.plot)

# Split data
set.seed(123)
index <- sample(1:nrow(bank_data), 0.7 * nrow(bank_data))
train <- bank_data[index, ]
test  <- bank_data[-index, ]
```

```
# Fit decision tree
dt_model <- rpart(y ~ ., data = train, method = "class", cp = 0.01)

printcp(dt_model)
```

Classification tree:
`rpart(formula = y ~ ., data = train, method = "class", cp = 0.01)`

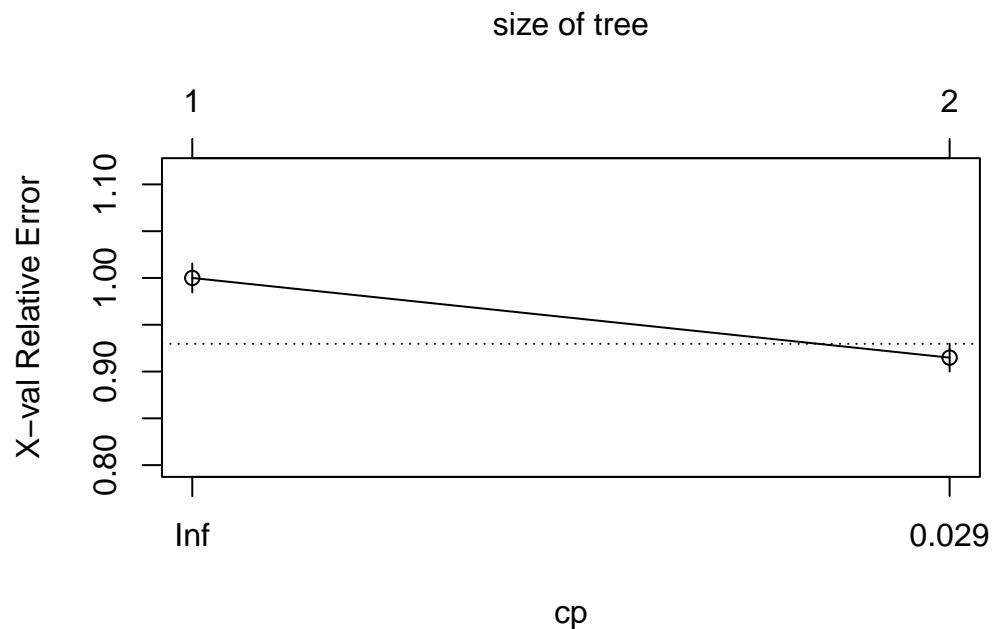
Variables actually used in tree construction:
`[1] poutcome`

Root node error: 3723/31647 = 0.11764

`n= 31647`

	CP	nsplit	rel error	xerror	xstd
1	0.085146	0	1.00000	1.00000	0.015395
2	0.010000	1	0.91485	0.91485	0.014808

```
plotcp(dt_model)
```

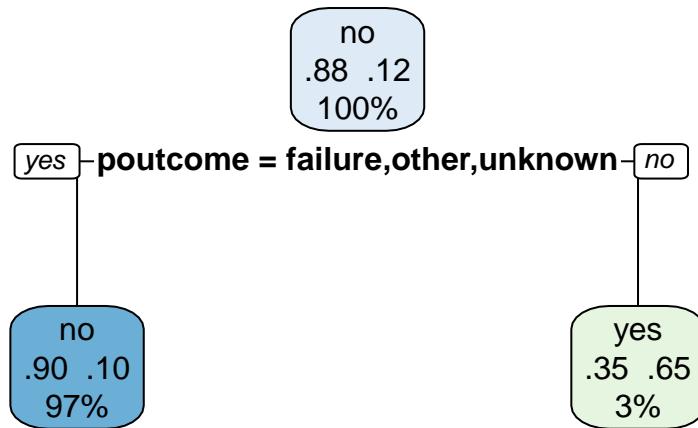


```

best_cp <- 0.029
dt_pruned <- prune(dt_model, cp = best_cp)

library(rpart.plot)
rpart.plot(dt_pruned, type = 2, extra = 104, fallen.leaves = TRUE)

```



```

dt_pred <- predict(dt_pruned, test, type = "class")
table(Predicted = dt_pred, Actual = test$y)

```

		Actual	
		Predicted	
		no	yes
no	11839	1279	
yes	159	287	

3. Decision Tree (Summary)

A pruned classification tree was built using cross-validation to determine the optimal complexity parameter ($cp = 0.029$). After pruning, the tree produced a single split based on *poutcome* (previous marketing outcome). This reflects that no other variable provided a sufficiently stable improvement in cross-validated accuracy.

The resulting model indicates that previous campaign success is the strongest discriminator: customers whose previous outcome was “success” show a markedly higher probability of subscribing (65%), whereas all other groups (“failure”, “other”, “unknown”) show very low subscription rates (10%). Although the tree has limited predictive performance due to the extreme class imbalance, it provides a clear and interpretable insight into which variable most differentiates customer behavior.

4. Cluster

```
# Data rangling
cluster_data <- train %>%
  select(age, balance, pdays, campaign, previous, poutcome)

cluster_dummy <- model.matrix(~ poutcome, data = cluster_data)[, -1]

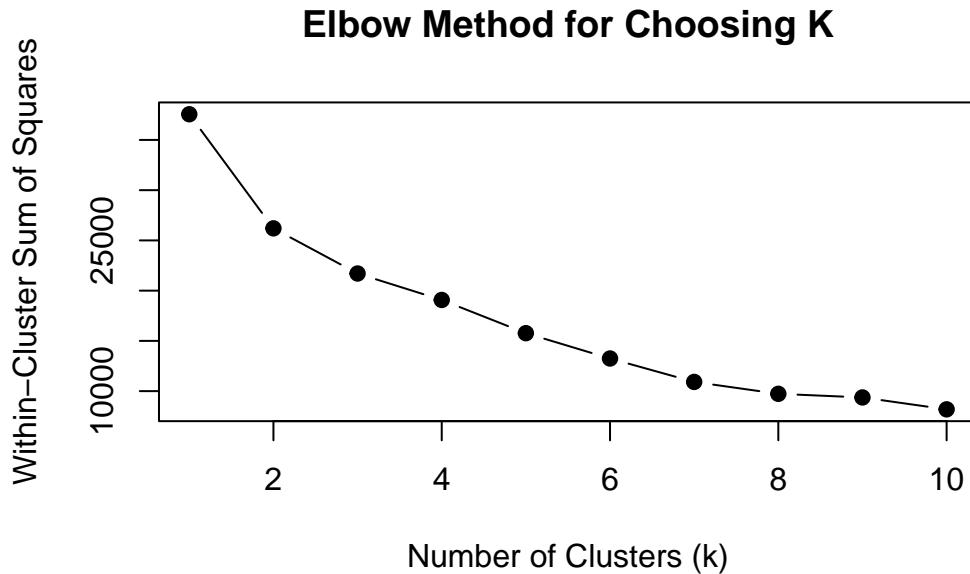
cluster_numeric <- cluster_data %>%
  select(-poutcome) %>%
  cbind(cluster_dummy)

cluster_scaled <- scale(cluster_numeric)

set.seed(123)
sample_rows <- sample(1:nrow(cluster_scaled), size = 0.15 * nrow(cluster_scaled))
cluster_sample <- cluster_scaled[sample_rows, ]

wss <- numeric(10)
for (k in 1:10) {
  km_temp <- kmeans(cluster_sample, centers = k, nstart = 20)
  wss[k] <- km_temp$tot.withinss
}

plot(1:10, wss, type = "b", pch = 19,
     xlab = "Number of Clusters (k)",
     ylab = "Within-Cluster Sum of Squares",
     main = "Elbow Method for Choosing K")
```



```
# run kmeans
k4 <- kmeans(cluster_scaled, centers = 4, nstart = 20)
```

```
# assign clusters to train
train$cluster <- k4$cluster
```

```
k4$centers
```

	age	balance	pdays	campaign	previous	poutcome	other
1	-0.097528878	0.02503037	1.848604	-0.09146813	1.3840437		4.8122015
2	0.198937561	0.22742820	1.222338	-0.30292825	1.0284279		-0.2077985
3	-0.001459108	-0.01514888	-0.412740	0.05113077	-0.2392114		-0.2077985
4	-0.013367815	0.03330833	1.991824	-0.25216554	0.9385283		-0.2077985

	poutcomesuccess	poutcomeunknown
1	-0.186610	-2.0999058
2	5.358599	-2.0999058
3	-0.186610	0.4761968
4	-0.186610	-2.0991647

```
train$cluster <- k4$cluster
```

```
cluster_yes_rate <- train %>%
```

```

group_by(k4$cluster) %>%
summarise(
  total = n(),
  yes = sum(y == "yes"),
  yes_rate = mean(y == "yes")
)
cluster_yes_rate

# A tibble: 4 x 4
`k4$cluster` total    yes yes_rate
<int> <int> <int>     <dbl>
1       1 1310    220  0.168
2       2 1065    691  0.649
3       3 25796   2383  0.0924
4       4 3476    429  0.123

train %>%
  group_by(cluster) %>%
  summarise(across(c(age, balance, pdays, campaign, previous), mean))

# A tibble: 4 x 6
cluster    age  balance   pdays campaign  previous
<int> <dbl>    <dbl>   <dbl>    <dbl>      <dbl>
1       1  39.9    1426.  228.     2.49  4.01
2       2  43.0    2034.  164.     1.82  3.13
3       3  40.9    1305. -0.996    2.94  0.0000388
4       4  40.7    1451.  242.     1.98  2.91

```

4. Cluster (Summary)

To identify distinct customer segments and uncover which groups are most likely to subscribe to the term deposit, we performed a K-means clustering analysis using key behavioral and demographic variables: age, account balance, days since last contact (pdays), number of campaign contacts, and previous contact history (previous and poutcome). After scaling the data and evaluating multiple values of k , a four-cluster solution was selected as the most interpretable and meaningful for marketing purposes.

The four clusters revealed clear differences in customer engagement patterns:

- **Cluster 3 – Warm Re-engagement Leads (Highest success rate: 64.9%)**

Customers in this group have been contacted multiple times in the past and show a moderate time gap since their last interaction. Their strong historical engagement makes them the most responsive segment and the primary target for focused follow-up campaigns.

- **Cluster 4 – Reactivation Candidates (Yes rate: 13.5%)**

Similar to Cluster 3 but with longer gaps since last contact. These customers still show above-average conversion potential and can be effectively reached through reactivation efforts.

- **Cluster 1 – Young, Low-Engagement Customers (Yes rate: 9.4%)**

This is the youngest segment with low balances and no prior contact history. Their low response rate suggests they are not ideal for high-cost outreach.

- **Cluster 2 – Older, Unfamiliar Customers (Yes rate: 9.0%)**

Although financially stable, this group has almost no previous contact experience. They tend to be less receptive to marketing messages and are not recommended as a priority segment.